

Project Proposal: Context-Aware Neural Speech Decoding with Skip-Diphone Auxiliary Supervision and Temporal Smoothness Regularization

Jiayu (Jarrod) Yang

Introduction to Deep Learning, Spring 2026 | Type-B Project (Domain Application)

1 Motivation

Speech brain–computer interfaces (BCIs) decode attempted speech from intracortical neural activity into text. The standard pipeline maps multi-channel neural features to frame-level phoneme probabilities with a recurrent neural network (RNN/GRU) under CTC loss, followed by a language model that converts phoneme predictions into words [1]. Predicting isolated (mono)phonemes is suboptimal because speech production is strongly context-dependent (*coarticulation*). The recently proposed DCoND framework [3] addresses this by predicting *diphones* ($z_{t-1} \rightarrow z_t$) and marginalizing over the previous-phoneme dimension to recover phoneme probabilities, achieving state-of-the-art Phoneme Error Rate (PER) on the Brain-to-Text '24 Benchmark.

This project builds on DCoND and tests two extensions. (i) **A skip-diphone auxiliary head** predicting the wider context pair $z_{t-2} \rightarrow z_t$ in addition to the standard adjacent diphone. The hypothesis is that speech motor planning encodes context beyond the immediately preceding phoneme; e.g., for *cat* (K AE T), the standard diphones cover K $\rightarrow$ AE and AE $\rightarrow$ T, while K $\rightarrow$ T probes whether earlier articulatory context still shapes later neural patterns. The auxiliary head is used only for supervision and is discarded at inference. (ii) **A temporal smoothness loss** on the marginalized phoneme probabilities. Because intracortical recordings are noisy and non-stationary, frame-level outputs can be jittery and trigger spurious phoneme insertions under CTC decoding; a smoothness regularizer encourages stable trajectories without preventing legitimate transitions.

2 Planned Data

I will use the **Brain-to-Text '24 Benchmark** (Eval.AI Challenge #2099 [4]), based on the speech BCI dataset of Willett et al. [2] (Dryad DOI 10.5061/dryad.x69p8czpq, Version 4). The dataset contains $\sim 12,100$ attempted-speech sentences from a single ALS participant (T12), recorded with intracortical microelectrode arrays in speech motor cortex; neural features are 20 ms-binned threshold-crossing rates and spike-band power. Specifically I will use: (a) `competitionData/` — official train/test/competitionHoldOut splits; (b) `languageModel/` — the released 3-gram (and optionally 5-gram) phoneme-to-text LM; (c) `derived/` — TFRecords and a baseline RNN checkpoint to validate the pipeline. Phoneme labels are obtained from sentence transcripts via the CMU Pronouncing Dictionary, following [1, 3]. My baseline codebase is the official PyTorch port `cffan/neural_seq_decoder` [5] by Fan (a co-author of [1]), which is the same codebase that DCoND [3] extends; the original TensorFlow release is `fwillett/speechBCI`. For debugging I will train on a small session subset, then scale to the full training set.

3 State-of-the-Art Approaches and Planned Selection

The NPTL baseline [1] is a bidirectional GRU mapping neural features to monophone probabilities under CTC; DCoND [3] replaces this target with a diphone target ($C^2=1600$ classes) and marginalizes back to phonemes. I follow DCoND's GRU encoder and CTC objective, and study the following ablation variants: (A) GRU + monophone CTC (NPTL baseline); (B) GRU + standard diphone with marginalization (DCoND); (C) B + temporal smoothness loss; (D) B + skip-diphone auxiliary head (predicting $z_{t-2} \rightarrow z_t$); (E) full model = B + skip-diphone aux. + smoothness. The training objective is

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{phoneme}}^{\text{CTC}} + \alpha \mathcal{L}_{\text{std-diphone}}^{\text{CTC}} + \beta \mathcal{L}_{\text{skip-diphone}}^{\text{CTC}} + \lambda \mathcal{L}_{\text{smooth}}, \quad \mathcal{L}_{\text{smooth}} = \frac{1}{T-1} \sum_{t=2}^T \|p_t - p_{t-1}\|_2^2, \quad (1)$$

where p_t is the marginalized phoneme distribution at frame t . Following [3] I initialize $\alpha=0.6$. I expect β and λ to be small so that auxiliary supervision and regularization *guide* — but do not dominate — the main CTC objective.

4 Evaluation Metrics

Primary metrics are **Phoneme Error Rate (PER)** and **Word Error Rate (WER)** after n -gram LM decoding, plus validation CTC loss for monitoring. I will report the five-variant ablation (A–E) above, sweep the smoothness weight $\lambda \in \{0, 10^{-3}, 5 \cdot 10^{-3}, 10^{-2}\}$, and run a small sweep over β . I hypothesize that moderate smoothness reduces frame-level jitter and improves PER/WER, while excessive smoothing erases phoneme boundaries; for the skip-diphone head, I will assess whether broader context helps under limited neural training data, or whether immediate diphones remain more reliable. Final-model results will also be submitted to the Eval.AI leaderboard [4] as an external sanity check.

Expected contribution. Beyond reproducing DCoND, this project tests whether (i) *broader phoneme context* via skip-diphone auxiliary supervision and (ii) *temporal output regularization* jointly improve neural speech decoding, yielding a context-aware, temporally regularized decoder for brain-to-text.

References [1] F. R. Willett *et al.*, A high-performance speech neuroprosthesis, *Nature* 620:1031–1036, 2023. [2] F. R. Willett *et al.*, Data for: A high-performance speech neuroprosthesis, *Dryad*, 2023, <https://doi.org/10.5061/dryad.x69p8czpq>. [3] J. Li, T. Le, C. Fan, M. Chen, E. Shlizerman, Brain-to-Text Decoding with Context-Aware Neural Representations and LLMs, *arXiv:2411.10657*, 2024. [4] Brain-to-Text Benchmark '24, Eval.AI Challenge #2099, <https://eval.ai/web/challenges/challenge-page/2099/overview>. [5] C. Fan *et al.*, Neural Sequence Decoder (PyTorch reference implementation), GitHub, https://github.com/cffan/neural_seq_decoder.